

# Variational Autoencoders on MNIST

## ELBO Derivation, KL Annealing, and CVAE Extension (v2)

Alexis Le Trung, Khatir Youyou, Tidjani Adam Kandine,  
Yahya Ahachim, Aniss Outaleb, Quentin Wurtlin

Generative AI and Diffusion Models

January 20, 2026

# Project Objectives

- ① **Derive the ELBO** from first principles with closed-form KL for Gaussians
- ② **Implement a Convolutional VAE** on MNIST with 2D latent space
- ③ **Track reconstruction vs KL** separately, implement KL annealing
- ④ **Visualize:** latent space, latent traversals, interpolations
- ⑤ **Extend to CVAE** for controlled digit generation

# ELBO Derivation

Starting from intractable marginal likelihood, we derived:

## Evidence Lower Bound

$$\log p(x) \geq \mathcal{L} = \underbrace{\mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)]}_{\text{Reconstruction (BCE)}} - \underbrace{\text{KL}(q_\phi(z|x)\|p(z))}_{\text{Regularization}}$$

**Closed-form KL** for Gaussian encoder  $q = \mathcal{N}(\mu, \sigma^2)$  and prior  $p = \mathcal{N}(0, I)$ :

$$\text{KL}(q\|p) = \frac{1}{2} \sum_{j=1}^d (\mu_j^2 + \sigma_j^2 - \log \sigma_j^2 - 1)$$

# ELBO — Proof of Lower Bound

Start from the marginal log-likelihood:

$$\log p_\theta(x) = \log \int p_\theta(x|z)p(z) dz$$

Introduce an arbitrary approximate posterior  $q_\phi(z|x)$  and rewrite:

$$\begin{aligned}\log p_\theta(x) &= \log \int q_\phi(z|x) \frac{p_\theta(x|z)p(z)}{q_\phi(z|x)} dz \\ &= \log \mathbb{E}_{q_\phi(z|x)} \left[ \frac{p_\theta(x|z)p(z)}{q_\phi(z|x)} \right]\end{aligned}$$

Apply Jensen's inequality ( $\log$  is concave):

$$\begin{aligned}\log p_\theta(x) &\geq \mathbb{E}_{q_\phi(z|x)} \left[ \log \frac{p_\theta(x|z)p(z)}{q_\phi(z|x)} \right] \\ &= \mathbb{E}_{q_\phi(z|x)} [\log p_\theta(x|z)] + \mathbb{E}_{q_\phi(z|x)} \left[ \log \frac{p(z)}{q_\phi(z|x)} \right] \\ &= \mathbb{E}_{q_\phi(z|x)} [\log p_\theta(x|z)] - \text{KL}(q_\phi(z|x) \| p(z))\end{aligned}$$

This establishes  $\log p_\theta(x) \geq \mathcal{L}(\theta, \phi; x)$ .

## ELBO — Gradients / Derivative

We optimize parameters  $\theta$  (decoder) and  $\phi$  (encoder) by maximizing the ELBO. The ELBO for a single datum is:

$$\mathcal{L}(\theta, \phi; x) = \mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)] - \text{KL}(q_\phi(z|x)\|p(z)).$$

Gradient w.r.t. decoder parameters  $\theta$  (decoder appears only in  $p_\theta(x|z)$ ):

$$\begin{aligned}\nabla_\theta \mathcal{L} &= \nabla_\theta \mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)] \\ &= \mathbb{E}_{q_\phi(z|x)} [\nabla_\theta \log p_\theta(x|z)].\end{aligned}$$

In practice estimate by Monte Carlo: sample  $z \sim q_\phi(z|x)$  via reparameterization  $z = \mu_\phi(x) + \sigma_\phi(x) \odot \epsilon$ .  
Gradient w.r.t. encoder parameters  $\phi$  uses reparameterization to push gradient through samples:

$$\nabla_\phi \mathcal{L} = \nabla_\phi \mathbb{E}_{\epsilon \sim \mathcal{N}(0, I)} [\log p_\theta(x|\mu_\phi + \sigma_\phi \odot \epsilon)] - \nabla_\phi \text{KL}(q_\phi\|p)$$

where the KL term has closed-form gradients:  $\partial_{\mu_j} \text{KL} = \mu_j$ ,  $\partial_{\sigma_j} \text{KL} = \sigma_j - 1/\sigma_j$ .

# Architecture & Training Setup

## Convolutional VAE:

- Encoder: Conv layers → 2D latent  $(\mu, \log \sigma^2)$
- Decoder: Linear → ConvTranspose layers
- Reparameterization:  $z = \mu + \sigma \odot \epsilon$


## KL Annealing:

$$\beta(t) = \min\left(1, \frac{t}{10}\right)$$


Prevents posterior collapse

Parameter	Value
Latent dim	2
Batch size	128
Learning rate	$10^{-3}$
Optimizer	Adam
Epochs	30
KL warmup	10 epochs

# VAE Training Curves



# Latent Space Visualization




## Observations:

- Clear clustering by digit class
- Similar digits nearby (4/9, 3/8, 1/7)
- Smooth, continuous manifold
- Matches  $\mathcal{N}(0, I)$  prior

## Key insight:

Organization emerges from reconstruction alone — no labels used!


# Latent Traversal Grid



Traversing  $z_1$  vs  $z_2$ :  $z_1$  (horizontal) controls slant/rotation;  $z_2$  (vertical) controls thickness/scale. 9 / 16

# Latent Interpolations

Latent Space Interpolations



Linear interpolation between encoded digit pairs. Smooth transitions validate the continuous, well-structured latent space.

# CVAE Extension: Conditional Generation

**Modification:** Condition encoder and decoder on class label  $c$

$$\mathcal{L}_{\text{CVAE}} = \mathbb{E}_{q_\phi(z|x,c)}[\log p_\theta(x|z,c)] - \text{KL}(q_\phi(z|x,c) \| p(z))$$

**Implementation:**

- Embedding the label to  $\times 28 \times 28$
- Concatenate with image as extra channels
- Class-independent prior:  $p(z) = \mathcal{N}(0, I)$

**Result:** Latent space encodes *style*, label provides *class identity*

# CVAE ELBO

Condition on label  $c$ ; objective per datum  $(x, c)$ :

$$\mathcal{L}_{\text{CVAE}}(\theta, \phi; x, c) = \mathbb{E}_{q_\phi(z|x, c)} [\log p_\theta(x|z, c)] - \text{KL}(q_\phi(z|x, c) \| p(z)).$$

Proof of lower bound follows identical steps with all densities conditioned on  $c$ :

$$\begin{aligned} \log p_\theta(x|c) &= \log \int p_\theta(x|z, c)p(z) dz \\ &= \log \mathbb{E}_{q_\phi(z|x, c)} \left[ \frac{p_\theta(x|z, c)p(z)}{q_\phi(z|x, c)} \right] \\ &\geq \mathbb{E}_{q_\phi(z|x, c)} \left[ \log \frac{p_\theta(x|z, c)p(z)}{q_\phi(z|x, c)} \right] = \mathcal{L}_{\text{CVAE}}. \end{aligned}$$

Gradients: similar decomposition as VAE; decoder gradient

$$\nabla_\theta \mathcal{L}_{\text{CVAE}} = \mathbb{E}_{q_\phi(z|x, c)} [\nabla_\theta \log p_\theta(x|z, c)]$$

Encoder gradients use reparameterization with the conditional encoder  $\mu_\phi(x, c), \sigma_\phi(x, c)$  and the closed-form KL.

## VAE vs CVAE: Quantitative Comparison

Model	Recon. Loss	KL	Total Loss
VAE	138.00	6.59	144.59
CVAE	124.36	4.80	129.16
<b>Improvement</b>	<b>9.9%</b>	<b>27.2%</b>	<b>10.7%</b>

### Why CVAE performs better:

- Decoder doesn't need to infer class from  $z \Rightarrow$  simpler task
- Latent space focuses purely on style variations
- More compact representation (lower KL)

# CVAE: Disentanglement Demonstration

CVAE: Disentangling Content ( $z$ ) and Class ( $c$ )



**Row 1-2:** Same  $z$  interpolation with different class labels (1 vs 7). **Row 3:** Fixed  $z$ , varying class 0-9  
⇒ style transfer across all digits.

# Key Findings

## ① ELBO provides principled training objective

- Reconstruction + regularization trade-off
- Closed-form KL enables efficient optimization

## ② KL annealing is essential

- Without it: posterior collapse,  $\text{KL} \rightarrow 0$
- 10-epoch warmup gave balanced training

## ③ VAE learns meaningful structure unsupervised

- Class clustering, smooth interpolations

## ④ CVAE achieves explicit disentanglement

- 9.9% better reconstruction, controlled generation

# Thank You!

Questions?

Code available in accompanying Jupyter notebook